

Chapter 2- Recommender Systems

Samuel Naser ¶

Basic techniques in recommendation

- Manhattan Distance (aka Taxi Cab Distance)
- Euclidean Distance

Example Distances

```
In [17]: users = {"Angelica": {"Blues Traveler": 3.5, "Broken Bells": 2.0,
    "Norah Jones": 4.5, "Phoenix": 5.0,
    "Slightly Stoopid": 1.5,
    "The Strokes": 2.5, "Vampire Weekend": 2.0},

    "Bill": {"Blues Traveler": 2.0, "Broken Bells": 3.5,
    "Deadmau5": 4.0, "Phoenix": 2.0,
    "Slightly Stoopid": 3.5, "Vampire Weekend": 3.0},

    "Chan": {"Blues Traveler": 5.0, "Broken Bells": 1.0,
    "Deadmau5": 1.0, "Norah Jones": 3.0,
    "Phoenix": 5, "Slightly Stoopid": 1.0},

    "Dan": {"Blues Traveler": 3.0, "Broken Bells": 4.0,
    "Deadmau5": 4.5, "Phoenix": 3.0,
    "Slightly Stoopid": 4.5, "The Strokes": 4.0,
    "Vampire Weekend": 2.0},

    "Hailey": {"Broken Bells": 4.0, "Deadmau5": 1.0,
    "Norah Jones": 4.0, "The Strokes": 4.0,
    "Vampire Weekend": 1.0},

    "Jordyn": {"Broken Bells": 4.5, "Deadmau5": 4.0, "Norah Jones": 5.0,
    "Phoenix": 5.0, "Slightly Stoopid": 4.5,
    "The Strokes": 4.0, "Vampire Weekend": 4.0},

    "Sam": {"Blues Traveler": 5.0, "Broken Bells": 2.0,
    "Norah Jones": 3.0, "Phoenix": 5.0,
    "Slightly Stoopid": 4.0, "The Strokes": 5.0},

    "Veronica": {"Blues Traveler": 3.0, "Norah Jones": 5.0,
    "Phoenix": 4.0, "Slightly Stoopid": 2.5,
    "The Strokes": 3.0}}
```

```
In [18]: #implementation for manhattan distance
def manhattanDistance(user1, user2):
    manhattanDistance = 0
    for band in user1:
        if band in user2:
            difference = user1[band] - user2[band]
            manhattanDistance += abs(difference)
    return manhattanDistance

#implementation for euclidean distance
def euclideanDistance(user1, user2):
    euclideanDistance = 0
    for band in user1:
        if band in user2:
            difference = user1[band] - user2[band]
            euclideanDistance += difference ** 2
    return round(euclideanDistance ** (0.5), 2)

print("Manhattan Distance between Angelica and Chan:",
manhattanDistance(users["Angelica"], users["Chan"]))
print("Euclidean Distance between Angelica and Chan:",
euclideanDistance(users["Angelica"], users["Chan"]))
```

Manhattan Distance between Angelica and Chan: 4.5

Euclidean Distance between Angelica and Chan: 2.4

NOW COMPUTE CLOSEST MATCH FOR USER!

```
In [35]: def nearestMatches(username, allUserData):
    allDistances = []
    for currentUser in allUserData:
        if currentUser != username:
            currentDistance = manhattanDistance(allUserData[username], a
llUserData[currentUser])
            allDistances.append((currentDistance, currentUser))
    allDistances.sort()
    return allDistances

print("All distances for user Veronica:", nearestMatches("Veronica", use
rs))
```

All distances for user Veronica: [(2.0, 'Hailey'), (3.5, 'Angelica'), (4.0, 'Bill'), (4.0, 'Dan'), (4.0, 'Jordyn'), (6.5, 'Chan'), (8.5, 'Sam')]

```
In [53]: def makeReccomendations(username, allUserData):
    recs = []
    nearestMatch = nearestMatches(username, allUserData)[0][1]
    nearestMatchRatings = allUserData[nearestMatch]
    userRatings = allUserData[username]
    for band in nearestMatchRatings:
        if not band in userRatings:
            recs.append((band, nearestMatchRatings[band]))
    return sorted(recs, key=lambda tup: tup[1], reverse=True) #don't fully understand this line

print(makeReccomendations("Veronica", users))

[('Broken Bells', 4.0), ('Deadmau5', 1.0), ('Vampire Weekend', 1.0)]
```

ALGORITHM, AS IT IS, WILL NOT RECOMMEND USERS BANDS IF THEIR CLOSEST MATCH HAS IS SUPERSET OF THEIR PREFERENCES

Assignment: Implement Minkowski Distance

```
In [55]: def minkowskiDistance(user1, user2, r):
    minkowskiDistance = 0
    for band in user1:
        if band in user2:
            difference = user1[band] - user2[band]
            minkowskiDistance += difference ** r
    return round(minkowskiDistance ** (1.0/r), 2)

print("Minkowski Distance between Angelica and Chan:",
minkowskiDistance(users["Angelica"], users["Chan"], 3))

Minkowski Distance between Angelica and Chan: 1.04
```

How to normalize user preferences?

Some users WAY more willing to give high ratings, other users are reluctant to drop a 5.

In data mining, we call this **GRADE INFLATION**

Pearson correlation coefficient to address this problem

Takes into account difference in user rating patterns.

```
In [7]: users = {"Angelica": {"Blues Traveler": 3.5, "Broken Bells": 2.0,
    "Norah Jones": 4.5, "Phoenix": 5.0,
    "Slightly Stoopid": 1.5,
    "The Strokes": 2.5, "Vampire Weekend": 2.0},

    "Bill": {"Blues Traveler": 2.0, "Broken Bells": 3.5,
    "Deadmau5": 4.0, "Phoenix": 2.0,
    "Slightly Stoopid": 3.5, "Vampire Weekend": 3.0},

    "Chan": {"Blues Traveler": 5.0, "Broken Bells": 1.0,
    "Deadmau5": 1.0, "Norah Jones": 3.0,
    "Phoenix": 5, "Slightly Stoopid": 1.0},

    "Dan": {"Blues Traveler": 3.0, "Broken Bells": 4.0,
    "Deadmau5": 4.5, "Phoenix": 3.0,
    "Slightly Stoopid": 4.5, "The Strokes": 4.0,
    "Vampire Weekend": 2.0},

    "Hailey": {"Broken Bells": 4.0, "Deadmau5": 1.0,
    "Norah Jones": 4.0, "The Strokes": 4.0,
    "Vampire Weekend": 1.0},

    "Jordyn": {"Broken Bells": 4.5, "Deadmau5": 4.0, "Norah Jones": 5.0,
    "Phoenix": 5.0, "Slightly Stoopid": 4.5,
    "The Strokes": 4.0, "Vampire Weekend": 4.0},

    "Sam": {"Blues Traveler": 5.0, "Broken Bells": 2.0,
    "Norah Jones": 3.0, "Phoenix": 5.0,
    "Slightly Stoopid": 4.0, "The Strokes": 5.0},

    "Veronica": {"Blues Traveler": 3.0, "Norah Jones": 5.0,
    "Phoenix": 4.0, "Slightly Stoopid": 2.5,
    "The Strokes": 3.0}}
```

```
In [16]: from math import sqrt
def pearsonCorrelation(r1, r2):
    sxy = 0
    sx = 0
    sy = 0
    sx2 = 0
    sy2 = 0
    n = 0
    for key in r1:
        if key in r2:
            n += 1
            x = r1[key]
            y = r2[key]
            sxy += x * y
            sx += x
            sy += y
            sx2 += x ** 2
            sy2 += y ** 2
    #if the denominator is zero, might as well return 0
    lower = sqrt(sx2 - sx ** 2 / n) * sqrt(sy2 - sy ** 2 / n)
    if lower == 0:
        return 0
    else:
        return (sxy - (sx * sy) / n) / lower

print("Pearson correlation between Angelican and Chan:", round(pearsonCorrelation(users["Angelica"], users["Chan"]),2))
```

Pearson correlation between Angelican and Chan: 0.82

Cosine Similarity

Used more frequently in text mining, but also used in collaborative filtering

In some cases, like text mining, shared zeroes do not indicate similarity (e.g. just because two things do **NOT** include word doesn't mean they are similar)

```
In [20]: #my cosine impl
def cosine(r1, r2):
    #calculate magnitudes of vectors
    m1 = sum([r1[key] ** 2 for key in r1]) ** 0.5
    m2 = sum([r2[key] ** 2 for key in r2]) ** 0.5
    #iterate over one vector checking for values for same key in other
    #NOTE: non-shared keys do not contribute to dot product
    dotProduct = 0
    for key in r1:
        if key in r2:
            dotProduct += r1[key] * r2[key]
    return dotProduct / (m1 * m2)

print(cosine(users["Angelica"], users["Chan"]))

0.8784261605942703
```

HUGE POINT: WHEN TO USE WHICH METHOD??

- With sparse data, where most values are zero, use Cosine similarity
- With dense data, where most data has values, and magnitude matters, use Manhattan or Euclidean distance
- With data subject to users rating on different scales, use a correlation value (Pearson coefficient)

Cannot just add zeroes to make sparse datasets work with Manhattan/Euclidean! The zeroes tend to dominate any calculations (you really just shouldn't take into account songs they haven't both rated).

- One possible workaround is to take a pseudo-average by finding the distance and then dividing by number of terms in common. This would, as you could see, result in lower distance (closer match) with *more* ratings in common

K-Nearest Neighbors

Takes out some of the quirks of basing recommendations of one (most) similar person

Use top k neighbors, and each neighbor influences the predictions by a factor of their similarity match to the user in question

```
In [99]: import csv
         from enum import Enum

         class Metric(Enum):
             PEARSON = 0
```



```

COSINE = 1
EUCLIDEAN = 2
MANHATTAN = 3

class Recommender:
    def __init__(self, _data, _metric):
        self.data = _data

        #assign appropriate correlation method
        if _metric == Metric.PEARSON:
            self.similarityMethod = self.__pearsonCorrelation
        if _metric == Metric.COSINE:
            self.similarityMethod = self.__cosine
        if _metric == Metric.EUCLIDEAN:
            self.similarityMethod = self.__euclidean
        if _metric == Metric.MANHATTAN:
            self.similarityMethod = self.__manhattan

    def recommend(self, user, amount):

        recommendations = {}
        #retrieve 3 nearest neighbors
        nearestUsers = self.kNearestNeighbors(user,3)
        userRatings = self.data[user]

        aggregateDistance = 0.0
        for user in nearestUsers:
            aggregateDistance += user[1]

        for user in nearestUsers:
            currentUsername = user[0]
            currentWeight = user[1] / aggregateDistance
            currentUserRatings = self.data[currentUsername]

            for band in currentUserRatings:
                if not band in userRatings:
                    if not band in recommendations:
                        recommendations[band] = currentUserRatings[band]
                * currentWeight
            else:
                recommendations[band] = recommendations[band] +
currentUserRatings[band] * currentWeight

        recommendations = list(recommendations.items())
        recommendations.sort(key=lambda x: x[1], reverse = True)

        return recommendations[:amount]

    def similarity(self, u1, u2):
        return self.similarityMethod(self.data[u1], self.data[u2])

    #method accepts username
    def kNearestNeighbors(self, user, k):
        similarities = []
        for currentUsername in self.data:
            if not currentUsername == user:
                similarities.append((currentUsername, self.similarity(us

```

```

er, currentUsername)))
    similarities = sorted(similarities, key = lambda x: x[1])
    return similarities[:k]

#definitions of similarity measures
#each takes two user rating dictionaries, and returns similarity
#lower = more similar

def __cosine(self, r1, r2):
    #calculate magnitudes of vectors
    m1 = sum([r1[key] ** 2 for key in r1]) ** 0.5
    m2 = sum([r2[key] ** 2 for key in r2]) ** 0.5
    #iterate over one vector checking for values for same key in other

    #NOTE: non-shared keys do not contribute to dot product
    dotProduct = 0
    for key in r1:
        if key in r2:
            dotProduct += r1[key] * r2[key]
    return dotProduct / (m1 * m2)

def __pearsonCorrelation(self, r1, r2):
    sxy = 0
    sx = 0
    sy = 0
    sx2 = 0
    sy2 = 0
    n = 0
    for key in r1:
        if key in r2:
            n += 1
            x = r1[key]
            y = r2[key]
            sxy += x * y
            sx += x
            sy += y
            sx2 += x ** 2
            sy2 += y ** 2
    #if the denominator is zero, might as well return 0
    lower = sqrt(sx2 - sx ** 2 / n) * sqrt(sy2 - sy ** 2 / n)
    if lower == 0:
        return 0
    else:
        return abs((sxy - (sx * sy) / n) / lower)

#implementation for manhattan distance
def __manhattanDistance(self, user1, user2):
    manhattanDistance = 0
    for band in user1:
        if band in user2:
            difference = user1[band] - user2[band]
            manhattanDistance += abs(difference)
    return manhattanDistance

#implementation for euclidean distance
def __euclideanDistance(self, user1, user2):
    euclideanDistance = 0

```

```

        for band in user1:
            if band in user2:
                difference = user1[band] - user2[band]
                euclideanDistance += difference ** 2
        return round(euclideanDistance ** (0.5), 2)

#main section
movie_file_rows = []
with open("movie_ratings.csv", 'r+') as movies:
    reader = csv.reader(movies)
    for rows in reader:
        movie_file_rows.append(rows)

#let's parse this data
userData = {}

movie_count = len(movie_file_rows) - 1
user_count = len(movie_file_rows[0]) - 1

#get list of usernames
usernames = []
for i in range(user_count):
    usernames.append(movie_file_rows[0][i+1])

#get list of movie names
movie_names = []
for i in range(movie_count):
    movie_names.append(movie_file_rows[i+1][0])

#set each entry to zero, at first
for username in usernames:
    userData[username] = {}

#fill matrix with data
dataMatrix = []
for i in range(movie_count):
    currentMovieRatings = []
    for j in range(user_count):
        currentMovieRatings.append(movie_file_rows[i+1][j+1])
    dataMatrix.append(currentMovieRatings)

for movieIndex in range(movie_count):
    currentMovie = movie_names[movieIndex]
    for userIndex in range(user_count):
        currentValue = dataMatrix[movieIndex][userIndex]
        if not currentValue == '':
            userData[usernames[userIndex]][movie_names[movieIndex]] = in
t(currentValue)

rcmd = Recommender(userData, Metric.PEARSON)
recommendation = rcmd.recommend("Bryan", 2)

print("RECOMMENDATIONS FOR BRYAN: ", recommendation)

```

```

RECOMMENDATIONS FOR BRYAN: [('Scarface', 4.3182870505659885), ('Blade
Runner', 3.4269113897024672)]

```